

DCTCP VS. SWIFT

DATA CENTER CONGESTION CONTROL COMPARISON

AUTHORS: YASMIN MITKAL , OMAR GARAH, AMIR WATTED
ADVISOR: ERAN TAVOR

ABSTRACT

Efficient congestion control in data centers is critical to achieving low latency and high throughput under shallow-buffer conditions. We compare two modern protocols—DCTCP (Data Center TCP), which uses ECN-based proportional reactions, and Swift, a new and evolving delay-based congestion controller via NS-3 simulations and packet-level analysis.

OBJECTIVE

Compare DCTCP and Swift as two data center congestion control protocols under realistic data center conditions.

PROJECT OVERVIEW

- Understand the core principles of DCTCP and Swift.
- Ramp up and configure the NS-3 simulator.
- Build a fat-tree topology in NS-3.
- Analyze the existing implementation of DCTCP in NS-3.
- Implement the Swift protocol within NS-3.
- Design and run simulations under varying load levels and traffic patterns.
- Compare DCTCP and Swift performance across different scenarios.

PROTOCOLS OVERVIEW

DCTCP: ecn-based

- Switches mark packets with ECN when $queue > threshold\ K$
- Receiver echoes ECN marks; sender tracks fraction α of marked packets
- Congestion window updated: $cwnd \leftarrow cwnd \times (1 - \alpha/2)$
- Retains standard TCP features: slow start, additive increase, loss recovery

$$(1) \quad \alpha \leftarrow (1 - g) \times \alpha + g \times F$$
$$(2) \quad cwnd \leftarrow cwnd \times (1 - \alpha/2)$$

SWIFT: delay-based


- RTT as Congestion Signal: reacts to end-to-end delay rather than ECN marks
- Topology-Based Target Delay: per-flow RTT goal that scales with path hops and active flow count (topology- and load-aware)
- ACK-Driven AIMD: if $RTT \leq target \rightarrow$ additive increase; if $RTT > target \rightarrow$ multiplicative decrease proportional to delay excess
- Hardware-Based Delay Breakdown: uses NIC timestamps to separate fabric vs. endpoint queuing delays
- Incast Pacing: supports sub-packet cwnd (<1) with precise pacing to absorb large bursts

Algorithm 1: SWIFT REACTION TO CONGESTION
Parameters: α : additive increment, β : multiplicative decrease constant, max_mdf : maximum multiplicative decrease factor
 $cwnd_prev \leftarrow cwnd$
 $bool\ can_decrease \leftarrow$ » Enforces MD once every RTT
 $(now - t_last_decrease \geq rtt)$
On Receiving ACK
 $retransmit_cnt \leftarrow 0$
 $target_delay \leftarrow TargetDelay()$ » See S3.5
if $delay < target_delay$ **then** » Additive Increase (AI)
 if $cwnd \geq 1$ **then**
 $cwnd \leftarrow cwnd + \frac{\alpha}{cwnd} \cdot num_acked$
 else
 $cwnd \leftarrow cwnd + \alpha \cdot num_acked$
else » Multiplicative Decrease (MD)
 if $can_decrease$ **then**
 $cwnd \leftarrow max(1 - \beta \cdot (\frac{delay - target_delay}{delay}), cwnd_prev)$
 Target Delay :
$$t = base_target + \#hops \times \bar{h} + max(0, min(\frac{\alpha}{\sqrt{f \cdot cwnd}} + \beta, fs_range))$$

On Retransmit Timeout
 $retransmit_cnt \leftarrow retransmit_cnt + 1$
if $retransmit_cnt \geq RETX_RESET_THRESHOLD$ **then**
 $cwnd \leftarrow min_cwnd$
else
 if $can_decrease$ **then**
 $cwnd \leftarrow (1 - max_mdf) \cdot cwnd$
On Fast Recovery
 $retransmit_cnt \leftarrow 0$
if $can_decrease$ **then**
 $cwnd \leftarrow (1 - max_mdf) \cdot cwnd$
 $cwnd \leftarrow$ » Enforce lower/upper bounds
 $clamp(min_cwnd, cwnd, max_cwnd)$
if $cwnd \leq cwnd_prev$ **then**
 $t_last_decrease \leftarrow now$
if $cwnd < 1$ **then**
 $pacing_delay \leftarrow \frac{rtt}{cwnd}$
else
 $pacing_delay \leftarrow 0;$

TARGET DELAY CONFIGURATION


To configure Swift's target delay, we experimented with various fixed delay values. We observed how each value impacted throughput, latency, and packet loss.



- Best Strategy: Using the minimum observed RTT (baseRTT) as the target delay achieved the best balance across all metrics.
- Key Insights:
 - Confirms Swift paper's trade-off:
 - Larger delays increase throughput but cause higher latency and packet loss.
 - Smaller delays reduce latency but limit throughput.
- Conclusion: Fine-tuning the target delay is crucial. Static delays must be chosen carefully to avoid under-utilizing bandwidth or overloading the network.


METHODOLOGY & EXPERIMENTAL SETUP

- Fat-tree topology in NS-3.
- Multiple traffic modes.
- Metrics: Throughput, Delay, Packet Loss.
- On-off application
- Averaged over 3 runs per data point.
- 3 swift variants:
 - Static target delay – 25ms
 - Target delay is base rtt (min rtt observed)
 - Target delay is Base rtt + $\epsilon \sim 1ms$



RESULTS

Results for K=12 (432 hosts) and a 7 seconds simulation runtime - random hot spot:



CONCLUSIONS

- Swift (baseRTT) and DCTCP showed similar throughput. Swift had lower packet loss, DCTCP maintained lower delay.
- Results match Swift's original paper claims on throughput and loss. However, delay improvements were not fully replicated in our setup.
- Performance is highly sensitive to parameters and target delay tuning.
- Target Delay Trade-off: Smaller delays reduce queuing and loss, larger ones boost throughput early but degrade performance under load. baseRTT provides a good balance.
- Overall Insight: Swift's delay-based approach is promising, requires refined tuning of delay targets and pacing.

ACHEIVEMENTS

- Implemented Swift in NS-3 and uploaded it to GitHub to support the community.
- Fine-tuned Swift parameters for improved performance and stability.
- Developed automation tools to streamline simulation and result analysis.
- Built a solid base for future projects involving Swift and congestion control protocols.
- Compared a new, developing protocol (Swift) against an industry-established protocol (DCTCP).

